

Import Modules

```
In [ ]: from luma.preprocessing.scaler import StandardScaler
        from luma.model_selection.split import TrainTestSplit
        from luma.neural.multi_layer import MLPClassifier
        from luma.visual.evaluation import ConfusionMatrix, ROCCurve, PrecisionRecallCurve

        from sklearn.datasets import fetch_openml
        import matplotlib.pyplot as plt
        import numpy as np
```

Load MNIST Dataset

```
In [ ]: X, y = fetch_openml('mnist_784', version=1, as_frame=False, return_X_y=True)
        X.shape, len(np.unique(y))
```

/Users/chanlee/miniforge3/lib/python3.10/site-packages/sklearn/datasets/_openml.py:1022: FutureWarning: The default value of 'parser' will change from 'liac-arff' to 'auto' in 1.4. You can set 'parser='auto'' to silence this warning. Therefore, an 'ImportError' will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.

```
Out[ ]: ((70000, 784), 10)
```

Split Dataset

```
In [ ]: X_train, X_test, y_train, y_test = TrainTestSplit(X, y.astype(int),
                                                         test_size=0.2,
                                                         random_state=42).get

        X_train.shape, X_test.shape
```

```
Out[ ]: ((56000, 784), (14000, 784))
```

Standardize Through StandardScaler

```
In [ ]: sc = StandardScaler()
        X_train_std = sc.fit_transform(X_train)
        X_test_std = sc.fit_transform(X_test)
```

Construct MLP

```
In [ ]: mlp = MLPClassifier(input_size=784,
                            hidden_sizes=128,
                            output_size=10,
                            max_epoch=1000,
                            learning_rate=1e-5,
                            lambda_=0.01,
                            dropout_rate=0.05,
                            activation='relu',
                            verbose=50)
```

Train MLP

```
In [ ]: mlp.fit(X_train_std, y_train)
        mlp.score(X_test_std, y_test)
```

```
[MLPClassifier] Epoch 0, Loss: 2.303604262984684
[MLPClassifier] Epoch 50, Loss: 0.2606702647666447
[MLPClassifier] Epoch 100, Loss: 0.19407239894880063
[MLPClassifier] Epoch 150, Loss: 0.15794643453507295
[MLPClassifier] Epoch 200, Loss: 0.13467669310660701
[MLPClassifier] Epoch 250, Loss: 0.11623951349414803
[MLPClassifier] Epoch 300, Loss: 0.10322330307000782
[MLPClassifier] Epoch 350, Loss: 0.09300160316609234
[MLPClassifier] Epoch 400, Loss: 0.08355726013321457
[MLPClassifier] Epoch 450, Loss: 0.07769183200887742
[MLPClassifier] Epoch 500, Loss: 0.07016529453290743
[MLPClassifier] Epoch 550, Loss: 0.06601742150669494
[MLPClassifier] Epoch 600, Loss: 0.061357056198030384
[MLPClassifier] Epoch 650, Loss: 0.0567581933343201
[MLPClassifier] Epoch 700, Loss: 0.05523374541465102
[MLPClassifier] Epoch 750, Loss: 0.050975375587188225
[MLPClassifier] Epoch 800, Loss: 0.048807523655764254
```

```
[MLPClassifier] Epoch 850, Loss: 0.047306548084005515
[MLPClassifier] Epoch 900, Loss: 0.04571841461863895
[MLPClassifier] Epoch 950, Loss: 0.04268030763919107
Out [ ]: 0.9696428571428571
```

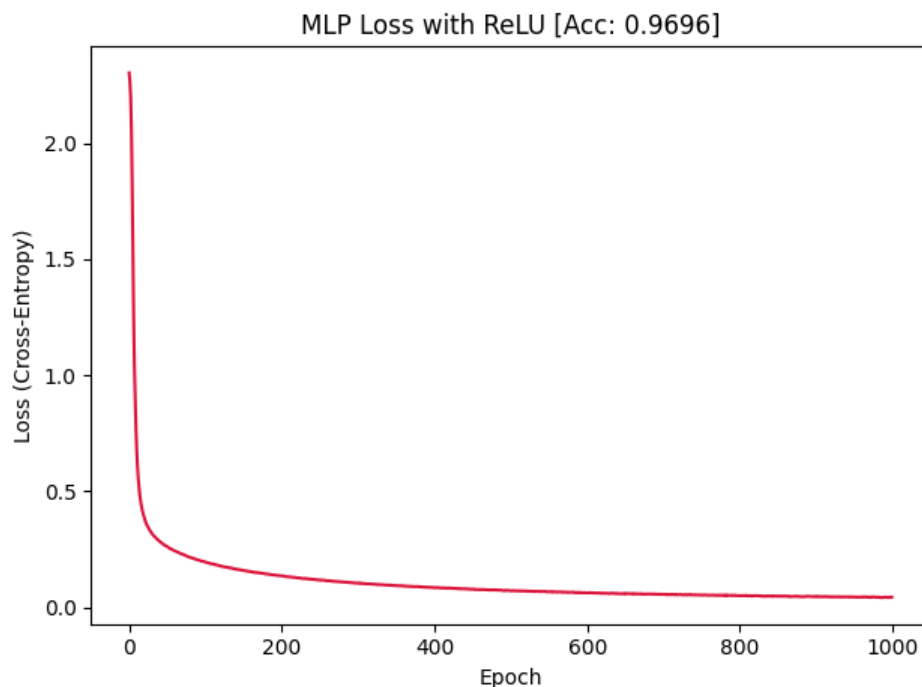
Print MLP Configuration

```
In [ ]: mlp.dump()
```

```
-----
|                               MLP Configuration                               |
|                               Input Size : 784                             |
| Layer 1 (Hidden) : 784 -> 128, Parameters: 100352 + 128 = 100480         |
| Layer 2 (Output) : 128 -> 10, Parameters: 1280 + 10 = 1290              |
|                               Total Parameters: 101770                     |
|                               Activation Function: ReLU                     |
|-----|
```

Plot Losses

```
In [ ]: plt.plot(range(mlp.max_epoch), mlp.losses_, c='crimson')
plt.xlabel('Epoch')
plt.ylabel('Loss (Cross-Entropy)')
plt.title(f'MLP Loss with {type(mlp.act_).__name__} [Acc: {mlp.score(X_test_std, y_test):.4f}]')
plt.tight_layout()
plt.show()
```

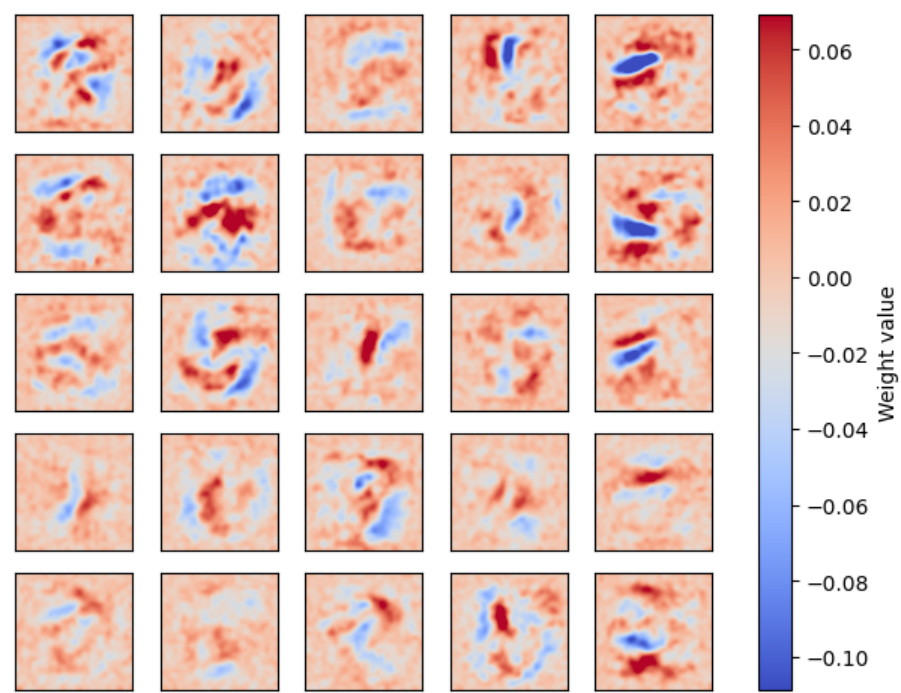


Visualize Weights

```
In [ ]: fig, axes = plt.subplots(5, 5)
global_vmin = 0.4 * mlp.weights[0].min()
global_vmax = 0.4 * mlp.weights[0].max()


for coef, ax in zip(mlp.weights[0].T, axes.ravel()):
    im = ax.matshow(coef.reshape(28, 28),
                    cmap='coolwarm',
                    vmin=global_vmin,
                    vmax=global_vmax,
                    interpolation='bicubic')
    ax.set_xticks(())
    ax.set_yticks(())

plt.tight_layout()
cbar = fig.colorbar(im, ax=axes.ravel().tolist())
cbar.set_label('Weight value')
plt.show()
```



ROC-AUC Curve


```
In [ ]: roc = ROCCurve(y_test, mlp.predict_proba(X_test_std))
        roc.plot(show=True)
```



```
Out [ ]: <Axes: title={'center': 'ROC Curve'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

PR Curve


```
In [ ]: pr = PrecisionRecallCurve(y_test, mlp.predict_proba(X_test_std))
        pr.plot(show=True)
```



Out[]: <Axes: title={'center': 'Precision-Recall Curve'}, xlabel='Recall', ylabel='Precision'>

Confusion Matrix

```
In [ ]: conf = ConfusionMatrix(y_test, mlp.predict(X_test_std))
conf.plot(show=True)
```



Out[]: <Axes: title={'center': 'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>

Visualize Sample Data Predictions

```
In [ ]: def plot_images(images, y_true, y_pred, cmap=plt.cm.gray_r):
fig, axes = plt.subplots(nrows=7, ncols=7, figsize=(7, 7))
fig.suptitle('Image Grid (True, Prediction)', fontsize=16)
axes = axes.flatten()
for img, ax, true, pred in zip(images, axes, y_true, y_pred):
ax.imshow(img, cmap=cmap, interpolation='nearest')
ax.set_title(f"{true}, {pred}, {true == pred}")
ax.axis('off')
plt.tight_layout()
plt.subplots_adjust(top=0.9)
plt.show()
```

```
In [ ]: plot_images(X_test[:49].reshape(-1, 28, 28), y_test[:49], mlp.predict(X_test_std[:49]))
```

